

Laplacian Paradigm 2.0

8:40-9:10: Merging Continuous and Discrete(Richard Peng)

9:10-9:50: Beyond Laplacian Solvers (Aaron Sidford)

9:50-10:30: Approximate Gaussian Elimination (Sushant Sachdeva)

10:30-11:00: coffee break

11:00-12:00: Analysis using matrix Martingales (Rasmus Kyng)

12:00-14:00 lunch

14:00-15:00 Graph Structure via Eliminations (Aaron Schild)

Website: bit.ly/laplacian2

Merging the Continuous and Discrete

Richard Peng





Oct 6, 2018

Outline

- **Graphs and Laplacians**
- Building Blocks
- Laplacian Paradigm 2.0

Large Networks

- Data mining: centrality, clustering...
- Image/video processing: segmentation, denoising ...
- Scientific computing: stress, fluids, waves...



julia




- \ (linear system solve)
- CVX (convex optimization)
- Eigenvector solvers

Graphs and Matrices

High performance computing: non-zeros \Leftrightarrow edges,
design / analyze matrix algorithms using graph theory

2	-1	-1
-1	1	0
-1	0	1




n vertices
m edges

n rows / columns
 $O(m)$ non-zeros

Graphs and Matrices

High performance computing: non-zeros \Leftrightarrow edges,
design / analyze matrix algorithms using graph theory

2	-1	-1
-1	1	0
-1	0	1



n vertices
m edges

n rows / columns
 $O(m)$ non-zeros

graph Laplacian matrix \mathbf{L}

- Diagonal: degrees
- Off-diagonal: -edge weights

Source of Laplacians

$$\begin{bmatrix} 2 & -1 & -1 \\ -1 & 1 & 0 \\ -1 & 0 & 1 \end{bmatrix}$$

graph Laplacian matrix \mathbf{L}

- Diagonal: degrees
- Off-diagonal: -edge weights

d-Regular graphs: $\mathbf{L} = d\mathbf{I} - \mathbf{A}$, \mathbf{A} : adjacency matrix

Source of Laplacians


$$\begin{bmatrix} 2 & -1 & -1 \\ -1 & 1 & 0 \\ -1 & 0 & 1 \end{bmatrix}$$

graph Laplacian matrix \mathbf{L}

- Diagonal: degrees
- Off-diagonal: -edge weights

d-Regular graphs: $\mathbf{L} = d\mathbf{I} - \mathbf{A}$, \mathbf{A} : adjacency matrix

$$\text{Graph cuts: } \mathbf{x}^T \mathbf{L} \mathbf{x} = \sum_{u \sim v} \mathbf{w}_{uv} (\mathbf{x}_u - \mathbf{x}_v)^2$$



\mathbf{x} indicator vector of cut \rightarrow weight of cut

Source of Laplacians


$$\begin{bmatrix} 2 & -1 & -1 \\ -1 & 1 & 0 \\ -1 & 0 & 1 \end{bmatrix}$$

graph Laplacian matrix \mathbf{L}

- Diagonal: degrees
- Off-diagonal: -edge weights

d-Regular graphs: $\mathbf{L} = d\mathbf{I} - \mathbf{A}$, \mathbf{A} : adjacency matrix

$$\text{Graph cuts: } \mathbf{x}^T \mathbf{L} \mathbf{x} = \sum_{u \sim v} \mathbf{w}_{uv} (\mathbf{x}_u - \mathbf{x}_v)^2$$



\mathbf{x} indicator vector of cut \rightarrow weight of cut

$$\mathbf{L} = \mathbf{B}^T \mathbf{W} \mathbf{B} \text{ where } \mathbf{B} \text{ is edge-vertex incidence matrix}$$

Origin of the Laplacian Paradigm

[Spielman Teng '04]

Input: graph Laplacian \mathbf{L}

vector \mathbf{b}

Output: vector \mathbf{x} s.t. $\mathbf{L}\mathbf{x} \approx \mathbf{b}$

Runtime: $O(m \log^{O(1)} n \log(1/\epsilon))$

Origin of the Laplacian Paradigm

[Spielman Teng '04]

Input: graph Laplacian \mathbf{L}

vector \mathbf{b}

Output: vector \mathbf{x} s.t. $\mathbf{L}\mathbf{x} \approx \mathbf{b}$

Runtime: $O(m \log^{O(1)} n \log(1/\epsilon))$

[Cohen-Kyng-Miller-Pachocki-P-Rao-Xu '14]: $\leq 1/2$

Origin of the Laplacian Paradigm

[Spielman Teng '04]

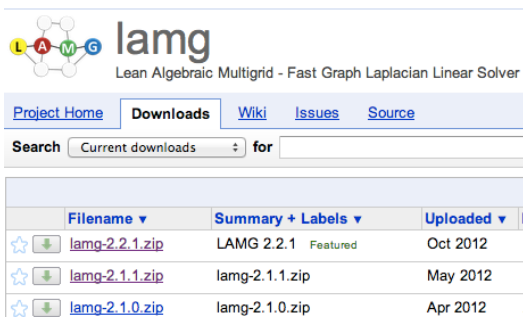
Input: graph Laplacian \mathbf{L}

vector \mathbf{b}

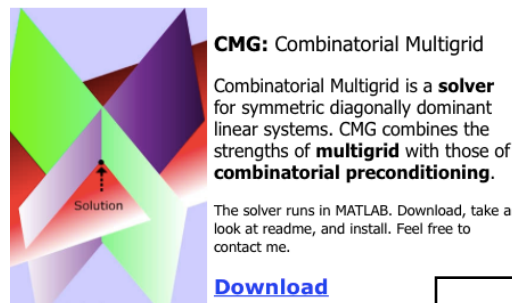
Output: vector \mathbf{x} s.t. $\mathbf{L}\mathbf{x} \approx \mathbf{b}$

Runtime: $O(m \log O(1) n \log(1/\epsilon))$

[Cohen-Kyng-Miller-Pachocki-P-Rao-Xu '14]: $\leq 1/2$

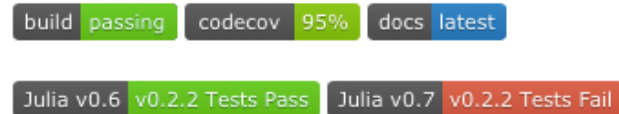


The screenshot shows the LAMG project website. At the top, there is a logo with the letters L, A, M, G in a circle, followed by the text "lamg" and "Lean Algebraic Multigrid - Fast Graph Laplacian Linear Solver". Below this, there are links for "Project Home", "Downloads", "Wiki", "Issues", and "Source". A search bar is present with the text "Current downloads" and "for". Below the search bar, there is a table with columns "Filename", "Summary + Labels", and "Uploaded". The table lists three versions of the software: "lamg-2.2.1.zip" (Oct 2012), "lamg-2.1.1.zip" (May 2012), and "lamg-2.1.0.zip" (Apr 2012).



The screenshot shows the CMG project website. At the top, there is a logo with the letters C, M, G in a circle, followed by the text "CMG: Combinatorial Multigrid". Below this, there is a description: "Combinatorial Multigrid is a solver for symmetric diagonally dominant linear systems. CMG combines the strengths of multigrid with those of combinatorial preconditioning." Below the description, there is a link for "Download".

Laplacians.jl




The screenshot shows the Laplacians.jl project website. At the top, there are links for "build", "passing", "codecov", "95%", "docs", and "latest". Below these links, there is a table with columns "Julia v0.6", "v0.2.2 Tests Pass", "Julia v0.7", and "v0.2.2 Tests Fail".

Wall clock: $m \leq 10^7$ in $\leq 20s$



The Laplacian Paradigm

Directly related:
Elliptic systems








$$Lx=b$$

Few iterations:
Eigenvectors,
Heat kernels



**Many iterations /
modify algorithm**
Graph problems
Image processing




Outline

- Graphs and Laplacians
- **Building Blocks**
- Laplacian Paradigm 2.0

$Lx = b$ as a graph problem

x : voltage vectors

Dual: electrical flow f



Unified formulation: \min_f with residual b $\|f\|_p$:

- $p = 2$: solving $Lx = b$
- $p = 1$: shortest path / transshipment
- $p = \infty$: max-flow/min-cut

The Carus Mathematical Monographs
NUMBER TWENTY-TWO

RANDOM WALKS
AND
ELECTRIC NETWORKS



PETER G. DOYLE
and
J. LAURIE SNELL


Direct Methods (combinatorial)

Repeatedly remove vertices by creating equivalent graphs on their neighborhoods

$$\mathbf{M}^{(2)} \leftarrow \text{Eliminate}(\mathbf{M}^{(1)}, i_1)$$

$$\mathbf{M}^{(3)} \leftarrow \text{Eliminate}(\mathbf{M}^{(2)}, i_2)$$

...



Direct Methods (combinatorial)


Repeatedly remove vertices by creating equivalent graphs on their neighborhoods

$$\mathbf{M}^{(2)} \leftarrow \text{Eliminate}(\mathbf{M}^{(1)}, i_1)$$

$$\mathbf{M}^{(3)} \leftarrow \text{Eliminate}(\mathbf{M}^{(2)}, i_2)$$


...

- Parallel graph algorithms
- Matrix multiplication / dense solves
- Sparsified squaring



Iterative Methods (numerical)

Solve $\mathbf{Ax} = \mathbf{b}$ by
 $\mathbf{x} \leftarrow \mathbf{x} - (\mathbf{Ax} - \mathbf{b})$




Fixed point: $\mathbf{Ax} - \mathbf{b} = 0$

Iterative Methods (numerical)

Preconditioning:

Solve $\mathbf{B}^{-1}\mathbf{Ax} = \mathbf{B}^{-1}\mathbf{b}$ by:

$$\mathbf{x} \leftarrow \mathbf{x} - \mathbf{B}^{-1}(\mathbf{Ax} - \mathbf{b})$$



Fixed point: $\mathbf{Ax} - \mathbf{b} = 0$


- Simple \mathbf{B} : $\mathbf{B} = \mathbf{I}$, many iterations
- $\mathbf{B} = \mathbf{A}$: 1 iteration, but same problem

Iterative Methods (numerical)

Preconditioning:

Solve $\mathbf{B}^{-1}\mathbf{Ax} = \mathbf{B}^{-1}\mathbf{b}$ by:

$$\mathbf{x} \leftarrow \mathbf{x} - \mathbf{B}^{-1}(\mathbf{Ax} - \mathbf{b})$$




Fixed point: $\mathbf{Ax} - \mathbf{b} = 0$

- Simple \mathbf{B} : $\mathbf{B} = \mathbf{I}$, many iterations
- $\mathbf{B} = \mathbf{A}$: 1 iteration, but same problem


- Krylov space methods / PCG
- Convex optimization algorithms

Hard instances

Direct methods create too much fill on highly connected graphs




Iterative methods take too many iterations paths




Hard instances

Direct methods create too much fill on highly connected graphs



Easy for iterative methods

Iterative methods take too many iterations paths




Easy for direct methods


Still 'easy' by themselves

Hard instances

Direct methods create too much fill on highly connected graphs



Iterative methods take too many iterations paths



Still 'easy' by themselves




Easy for iterative methods

Easy for direct methods

Must handle both simultaneously, but
avoid paying n iterations \times m per iteration






Hybrid algorithms (aka. v1.0)

- Scientific computing: iChol, multigrid
- [Vaidya '89] precondition with graphs







Hybrid algorithms (aka. v1.0)

- Scientific computing: iChol, multigrid
- [Vaidya '89] precondition with graphs



Hybrid algorithms (aka. v1.0)

- Scientific computing: iChol, multigrid
- [Vaidya '89] precondition with graphs




Focus: how to combine

- [Gemban-Miller '96]: spectral graph theory
- [Spielman-Teng '04]: spectral (ultra-)sparsify


Key “glue”: sparsification

[Spielman-Teng '04]: for any G , can find H with $O(n \log^{O(1)} n)$ edges s.t. $\mathbf{x}^T \mathbf{L}_G \mathbf{x} \approx \mathbf{x}^T \mathbf{L}_H \mathbf{x} \quad \forall \mathbf{x}$



Key “glue”: sparsification


[Spielman-Teng '04]: for any G , can find H with $O(n \log^{O(1)} n)$ edges s.t. $\mathbf{x}^T \mathbf{L}_G \mathbf{x} \approx \mathbf{x}^T \mathbf{L}_H \mathbf{x} \quad \forall \mathbf{x}$



- Combinatorial parameter: #edges
- Numerical parameter : approximations

Key “glue”: sparsification

[Spielman-Teng `04]: for any G , can find H with $O(n \log^{O(1)} n)$ edges s.t. $\mathbf{x}^T \mathbf{L}_G \mathbf{x} \approx \mathbf{x}^T \mathbf{L}_H \mathbf{x} \quad \forall \mathbf{x}$



[Spielman-Srivatava `08]: sample by effective resistances gives H with $O(n \log n)$ edges

Outline


- Graphs and Laplacians
- Building Blocks
- **Laplacian Paradigm 2.0**

Max-Flow problem

Maximum number of
disjoint s-t paths

Applications:

- Routing
- Scheduling



Recall: $\min_{\mathbf{f} \text{ with residual } \mathbf{b}} \|\mathbf{f}\|_p$:

- $p = 2$: solving $\mathbf{L}\mathbf{x} = \mathbf{b}$
- $p = \infty$: max-flow/min-cut


Max-Flow problem

Maximum number of disjoint s-t paths

Dual: separate s and t by removing fewest edges

Applications:

- Routing
- Scheduling



Applications:

- Partitioning
- Clustering


Recall: $\min_{\mathbf{f} \text{ with residual } \mathbf{b}} \|\mathbf{f}\|_p$:

- $p = 2$: solving $\mathbf{L}\mathbf{x} = \mathbf{b}$
- $p = \infty$: max-flow/min-cut

Hybrid Algorithms for Max-Flow

[Daitch-Spielman '08][Christiano-Kelner-Madry-Spielman-Teng '10]:

[Lee-Sidford '14] Max-flow/Min-cut via (several) electrical flows



Repeat about $m^{1/3}$ iters


- Solve linear systems
- Re-adjust edge weights

$$Lx=b$$

Hybrid Algorithms for Max-Flow

[Daitch-Spielman '08][Christiano-Kelner-Madry-Spielman-Teng '10]:

[Lee-Sidford '14] Max-flow/Min-cut via (several) electrical flows



Repeat about $m^{1/3}$ iters

- Solve linear systems
- Re-adjust edge weights


$$Lx=b$$

[Madry '10] [Racke-Shah-Taubig '14]:
cut approximator / oblivious routing
 $O(n^{o(1)})$ -approx. in $O(m^{1+o(1)})$

Hybrid Algorithms for Max-Flow

[Daitch-Spielman '08][Christiano-Kelner-Madry-Spielman-Teng '10]:

[Lee-Sidford '14] Max-flow/Min-cut via (several) electrical flows



[Lee-Rao-Srivastava '13][Sherman '13, '17][Kelner-Lee-Orecchia-Sidford '14]:
Preconditioning, $(1+\epsilon)$ -approx

Repeat about $m^{1/3}$ iters

- Solve linear systems
- Re-adjust edge weights


$$Lx=b$$

[Madry '10] [Racke-Shah-Taubig '14]:
cut approximator / oblivious routing
 $O(n^{o(1)})$ -approx. in $O(m^{1+o(1)})$

Hybrid Algorithms for Max-Flow

[Daitch-Spielman '08][Christiano-Kelner-Madry-Spielman-Teng '10]:

[Lee-Sidford '14] Max-flow/Min-cut via (several) electrical flows



Repeat about $m^{1/3}$ iters

- Solve linear systems
- Re-adjust edge weights


$$Lx=b$$

[Lee-Rao-Srivastava '13][Sherman '13, '17][Kelner-Lee-Orecchia-Sidford '14]:
Preconditioning, $(1+\epsilon)$ -approx



[P'16]: recurse them into each other:
 $O(m \log^{41} n)$, optimistically $m \log^6 n$

[Madry '10] [Racke-Shah-Taubig '14]:
cut approximator / oblivious routing
 $O(n^{o(1)})$ -approx. in $O(m^{1+o(1)})$

Laplacian Paradigm 2.0




Motivated by the goal of hybrid algorithms, modify direct and iterative methods




New Intermediate structures / theorems motivated by the overall algorithms

Examples

Directed graphs / asymmetric matrices




Sparsified/Approximate Gaussian Elimination




Under the hood

Matrix (martingale) concentration





Partitioning / Localizations of Random Walks




Not covered ☹️



Matrix Zoo from Scientific Computing



[Boman-Hendrickson-Vavasis `04]
[Kyng-Lee-P-Sachdeva-Spielman `16]
[Kyng-Zhang `17][Kyng-P-Schweiteman-Zhang `18]






Interactions with data structures



[Kelner-Orecchia-Sidford-Zhu `13]
[Nanongkai-Saranuk `17][Wulff-Nilsen `17]
[Durfee-Kyng-Peebles-Rao-Sachdeva `17]

Questions

Wist list	Direct	Iterative	Hybrid
Convex functions	?	☺	☺☺??
Arbitrary values	☺	?	?☹☺!
Dynamic/streaming	☺	☹	☺???



Approximate eliminations
beyond spectral condition #


Unreasonable effectiveness
of `pcg(ichol(A), b)`, multigrid

Non-linear (preconditioned)
iterative methods

[Adil-Kyng-P-Sachdeva `19]:
p-norm iterative refinement

Solvers in Practice

[Kyng-Rao-Sachdeva '15] we suggest rerunning the program a few times... An alternate solver based on iChol is provided...



Questions:

- Precision
- (pseudo) deterministic